


Grid carbon-intensity forecasting advantage: family-level vs. single-statistic



Dashed line = no-grid-data baseline. Panel (b): the mean alone (green) beats every other single statistic, including dispersion, as a STANDALONE predictor - this qualifies, but does not overturn, the SHAP finding that dispersion carries the largest attribution SHARE inside the full joint model (Fig. 7a): SHAP reports marginal contribution given the other grid features present, not standalone predictive power. The deployable conclusion is the FULL grid feature set, not any one statistic, drives the gain.